

## 03\_trends\_seasonality\_and\_forecasting.ipynb

```

%%
import os
from IPython.display import display
import matplotlib.ticker as mticker
import matplotlib.dates as mdates
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_absolute_error, mean_squared_error
import numpy as np
from sklearn.linear_model import Ridge
from pathlib import Path
import pandas as pd
import matplotlib.pyplot as plt

%%
repo_root = Path.cwd()
data_dir = (repo_root / "data") if (repo_root / "data").exists() else repo_root

paths = {
    'sales': data_dir / 'sales.csv',
    'stores': data_dir / 'stores.csv',
    'categories': data_dir / 'categories.csv',
    'dates': data_dir / 'dates.csv',
}

for k, p in paths.items():
    print(f"{k}: {p.exists()} -> {p.resolve()}")

def read_csv_smart(p):
    df = pd.read_csv(p, low_memory=False)
    df.columns = df.columns.str.lower().str.strip()
    date_cols = [c for c in df.columns if 'date' in c]
    for c in date_cols:
        df[c] = pd.to_datetime(df[c], errors='coerce')
    return df

sales_df = read_csv_smart(paths['sales'])
stores_df = read_csv_smart(paths['stores'])
categories_df = read_csv_smart(paths['categories'])
dates_df = read_csv_smart(paths['dates'])

try:
    import holidays
    years = range(int(dates_df['date'].dt.year.min()), int(dates_df['date'].dt.year.max()) + 1)
    us_holidays = holidays.US(years=years)
    holiday_dates = pd.to_datetime(list(us_holidays.keys()))
    dates_df['is_holiday'] = dates_df['date'].isin(holiday_dates).astype(int)
    print("\nHoliday flag on dates_df (0=non-holiday, 1=holiday):")
    print(dates_df['is_holiday'].value_counts().sort_index())
except Exception as e:
    print("\n[warn] Could not add holiday flag (install `holidays`?):", e)
    dates_df['is_holiday'] = 0

try:
    from meteostat import Daily, Point
    import numpy as np

    start_dt = pd.to_datetime(dates_df['date'].min()).to_pydatetime()
    end_dt = pd.to_datetime(dates_df['date'].max()).to_pydatetime()

    rep_point = None
    lat_cols = [c for c in stores_df.columns if c.lower() in ('lat', 'latitude')]
    lon_cols = [c for c in stores_df.columns if c.lower() in ('lon', 'lng', 'longitude')]

    if lat_cols and lon_cols:
        lat_med = pd.to_numeric(stores_df[lat_cols[0]], errors='coerce').median()
        lon_med = pd.to_numeric(stores_df[lon_cols[0]], errors='coerce').median()
        if pd.notna(lat_med) and pd.notna(lon_med):
            rep_point = Point(float(lat_med), float(lon_med))
            print(f"[weather] Using median store coords: lat={lat_med:.4f}, lon={lon_med:.4f}")

    if rep_point is None:
        rep_point = Point(39.8283, -98.5795)
        print("[weather] Using US centroid as representative location.")

```

```

wx = Daily(rep_point, start_dt, end_dt)
wx_full = wx.fetch()
wx_df = (
    wx_full.reset_index()
        .rename(columns={'time': 'date'})
        .loc[:, ['date', 'tavg', 'tmin', 'tmax', 'prcp']]
)
wx_df['date'] = pd.to_datetime(wx_df['date'])

if 'tavg' not in wx_df or wx_df['tavg'].isna().all():
    if {'tmin', 'tmax'}.issubset(wx_df.columns):
        wx_df['tavg'] = wx_df[['tmin', 'tmax']].mean(axis=1)
    else:
        doy = wx_df['date'].dt.dayofyear
        wx_df['tavg'] = (np.sin(2*np.pi*(doy/365.25)) * 10 + 13).round(1)
if 'prcp' not in wx_df:
    wx_df['prcp'] = 0.0

wx_df = wx_df[['date', 'tavg', 'prcp']]

dates_df = dates_df.merge(wx_df, on='date', how='left')

dates_df['tavg'] = dates_df['tavg'].ffill().bfill()
dates_df['prcp'] = dates_df['prcp'].fillna(0.0)

print("\nWeather columns added to dates_df (head):")
print(dates_df[['date', 'tavg', 'prcp']].head())

except Exception as e:
    print("[warn] Could not add weather features; proceeding without. Error:", e)
    try:
        import numpy as np
    except Exception:
        pass
    if 'tavg' not in dates_df.columns:
        if 'np' in globals():
            doy = dates_df['date'].dt.dayofyear
            dates_df['tavg'] = (np.sin(2*np.pi*(doy/365.25)) * 20 + 55).round(1)
        else:
            dates_df['tavg'] = pd.NA
    if 'prcp' not in dates_df.columns:
        dates_df['prcp'] = 0.0

for name, df in [('sales', sales_df), ('stores', stores_df),
                 ('categories', categories_df), ('dates', dates_df)]:
    print(f"\n{name}: {df.shape}")
    display(df.head(3))
    display(df.dtypes.to_frame('dtype').T)

###
sales2_df = (sales_df
             .merge(dates_df, on='date', how='left')
             .merge(stores_df, on='store_id', how='left')
             .merge(categories_df, on='category_id', how='left'))

sales2_df['weekend'] = sales2_df['dow'].isin([5,6]).astype(int)

print('shape:', sales2_df.shape)
print(sales2_df.columns.tolist())

print('\ndow counts:')
print(sales2_df['dow'].value_counts().sort_index())

print('\nweekend counts:')
print(sales2_df['weekend'].value_counts().sort_index())

daily_rev_df = (sales2_df.groupby('date')['sales_revenue']
                .sum().reset_index())
display(daily_rev_df.head())

###
sales = pd.read_csv("sales.csv", parse_dates=["date"])
cats = pd.read_csv("categories.csv")
df = sales.merge(cats, on="category_id")
df["doy"] = df["date"].dt.dayofyear
print("Category peak DOY:\n",

```

```

df.groupby(["name", "doy"])[ "sales_revenue"].mean()
    .reset_index()
    .loc[lambda x: x.groupby("name")[ "sales_revenue"].idxmax(), ["name", "doy"]]
    .sort_values("name"))

#%%
daily_rev_df = sales2_df.groupby('date', as_index=False)[ 'sales_revenue'].sum()

plt.figure(figsize=(12,4))
plt.plot(daily_rev_df['date'], daily_rev_df[ 'sales_revenue'])
plt.title('daily revenue (all stores)')
plt.xlabel('date'); plt.ylabel('revenue')
plt.tight_layout()
plt.show()

dow_avg_df = (sales2_df.groupby('dow')[ 'sales_revenue']
               .sum().div(sales2_df.groupby('dow')[ 'sales_revenue'].count())
               .reset_index(name='avg_rev'))
plt.figure(figsize=(6,4))
plt.bar(dow_avg_df['dow'], dow_avg_df[ 'avg_rev'])
plt.title('avg revenue by day-of-week (0=Mon)')
plt.xlabel('dow'); plt.ylabel('avg revenue')
plt.tight_layout()
plt.show()

wk_df = (sales2_df.groupby(['weekend', 'date'])[ 'sales_revenue']
         .sum().reset_index())
wk_pvt_df = wk_df.pivot(index='date', columns='weekend', values='sales_revenue').rename(columns={0: 'weekday', 1: 'weekend'})

plt.figure(figsize=(12,4))
plt.plot(wk_pvt_df.index, wk_pvt_df[ 'weekday'], label='weekday')
plt.plot(wk_pvt_df.index, wk_pvt_df[ 'weekend'], label='weekend', alpha=0.8)
plt.title('weekday vs weekend revenue')
plt.xlabel('date'); plt.ylabel('revenue'); plt.legend()
plt.tight_layout()
plt.show()

#%%
daily_rev_df = (sales2_df
               .assign(date=pd.to_datetime(sales2_df[ 'date']))
               .groupby('date', as_index=False)[ 'sales_revenue']
               .sum()
               .sort_values('date'))

daily_rev_df[ 'rev_30dma'] = (daily_rev_df[ 'sales_revenue']
                             .shift(1) # no peeking
                             .rolling(window=30, min_periods=7) # require some history
                             .mean())

plt.style.use("dark_background")
fig, ax = plt.subplots(figsize=(12, 4))

ax.plot(daily_rev_df[ 'date'], daily_rev_df[ 'sales_revenue'], alpha=0.8, label='daily')
ax.plot(daily_rev_df[ 'date'], daily_rev_df[ 'rev_30dma'], linewidth=2.2, label='30-day avg')

ax.set_title('daily revenue (all stores)')
ax.set_xlabel('date')
ax.set_ylabel('revenue')

ax.xaxis.set_minor_locator(mdates.MonthLocator(interval=3))
ax.xaxis.set_major_locator(mdates.YearLocator())
ax.xaxis.set_major_formatter(mdates.DateFormatter('%Y'))

ax.yaxis.set_major_locator(mticker.MultipleLocator(1000))

ax.grid(True, which='major', linestyle='--', color='gray', alpha=0.4)
ax.grid(True, which='minor', linestyle=':', color='gray', alpha=0.2)

ax.legend()
plt.tight_layout()
plt.show()

#%%
wk_df = (
    sales2_df.groupby(['weekend', 'date'])[ 'sales_revenue']
    .sum()

```

```

    .reset_index()
)

wk_pvt_df = (
    wk_df.pivot(index='date', columns='weekend', values='sales_revenue')
    .rename(columns={0: 'weekday', 1: 'weekend'})
)

plt.style.use("dark_background")
fig, ax = plt.subplots(figsize=(12, 4))

ax.plot(wk_pvt_df.index, wk_pvt_df['weekday'], label='weekday', linewidth=1.5)
ax.plot(wk_pvt_df.index, wk_pvt_df['weekend'], label='weekend', linewidth=1.5, alpha=0.8)

ax.set_title('weekday vs weekend revenue')
ax.set_xlabel('date')
ax.set_ylabel('revenue')

ax.xaxis.set_minor_locator(mdates.MonthLocator(interval=3))

ax.xaxis.set_major_locator(mdates.YearLocator())
ax.xaxis.set_major_formatter(mdates.DateFormatter('%Y'))

ax.yaxis.set_major_locator(plt.MultipleLocator(1000))

ax.grid(True, which='major', linestyle='--', color='gray', alpha=0.4)
ax.grid(True, which='minor', linestyle=':', color='gray', alpha=0.2)

ax.legend()
plt.tight_layout()
plt.show()

#%%
monthly_df = (sales2_df.groupby(['year', 'month'])['sales_revenue']
               .sum().reset_index())
monthly_df['date'] = pd.to_datetime(monthly_df['year'].astype(str) + '-' + monthly_df['month'].astype(str) + '-01')

plt.figure(figsize=(12,4))
plt.plot(monthly_df['date'], monthly_df['sales_revenue'])
plt.title('monthly revenue (all stores)')
plt.xlabel('month'); plt.ylabel('revenue')
plt.tight_layout()
plt.show()

cat_df = (sales2_df.groupby(['category_id', 'year', 'month'])['sales_revenue']
          .sum().reset_index())
cat_df['date'] = pd.to_datetime(cat_df['year'].astype(str) + '-' + cat_df['month'].astype(str) + '-01')
cat_pvt_df = cat_df.pivot(index='date', columns='category_id', values='sales_revenue')

cat_pvt_df.plot(figsize=(12,4))
plt.title('monthly revenue by category')
plt.xlabel('month'); plt.ylabel('revenue')
plt.legend(title='category_id')
plt.tight_layout()
plt.show()

#%%
monthly_df = (
    sales2_df.groupby(['year', 'month'])['sales_revenue']
    .sum()
    .reset_index()
)
monthly_df['date'] = pd.to_datetime(
    monthly_df['year'].astype(str) + '-' + monthly_df['month'].astype(str) + '-01'
)

plt.style.use("dark_background")
plt.figure(figsize=(12,4))
plt.plot(monthly_df['date'], monthly_df['sales_revenue'], linewidth=1.8)
plt.title('monthly revenue (all stores)')
plt.xlabel('month')
plt.ylabel('revenue')
plt.tight_layout()
plt.show()

cat_df = (
    sales2_df.groupby(['category_id', 'year', 'month'])['sales_revenue']

```

```

        .sum()
        .reset_index()
    )
    cat_df['date'] = pd.to_datetime(
        cat_df['year'].astype(str) + '-' + cat_df['month'].astype(str) + '-01'
    )

    cat_pvt_df = cat_df.pivot(index='date', columns='category_id', values='sales_revenue')

    category_names = {1: 'Beverages', 2: 'Bakery', 3: 'Meat', 4: 'Produce'}
    cat_pvt_df = cat_pvt_df.rename(columns=category_names)

    cat_pvt_df.plot(figsize=(12,4), linewidth=1.5)
    plt.title('monthly revenue by category')
    plt.xlabel('month')
    plt.ylabel('revenue')
    plt.legend(title='category')
    plt.tight_layout()
    plt.show()

    ###
    store_m_df = (sales2_df.groupby(['store_id', 'year', 'month'])['sales_revenue']
        .sum().reset_index())
    store_m_df['date'] = pd.to_datetime(store_m_df['year'].astype(str) + '-' + store_m_df['month'].astype(str) + '-01')
    store_pvt_df = store_m_df.pivot(index='date', columns='store_id', values='sales_revenue')

    store_pvt_df.plot(figsize=(12,4))
    plt.title('monthly revenue by store')
    plt.xlabel('month'); plt.ylabel('revenue')
    plt.legend(title='store_id'); plt.tight_layout(); plt.show()

    reg_m_df = (sales2_df.groupby(['region', 'year', 'month'])['sales_revenue']
        .sum().reset_index())
    reg_m_df['date'] = pd.to_datetime(reg_m_df['year'].astype(str) + '-' + reg_m_df['month'].astype(str) + '-01')
    reg_pvt_df = reg_m_df.pivot(index='date', columns='region', values='sales_revenue')

    reg_pvt_df.plot(figsize=(12,4))
    plt.title('monthly revenue by region')
    plt.xlabel('month'); plt.ylabel('revenue')
    plt.legend(title='region'); plt.tight_layout(); plt.show()

    ###
    ts_df = (sales2_df.groupby('date', as_index=False)['sales_revenue']
        .sum().rename(columns={'sales_revenue': 'y'}))

    full_idx = pd.date_range(ts_df['date'].min(), ts_df['date'].max(), freq='D')
    ts_df = (pd.DataFrame({'date': full_idx})
        .merge(ts_df, on='date', how='left')
        .fillna({'y': 0}))

    feat_cols = ['date', 'dow', 'woy', 'month', 'quarter', 'year', 'is_holiday', 'tavg', 'prcp']
    ts_df = ts_df.merge(dates_df[feat_cols], on='date', how='left')
    ts_df['weekend'] = ts_df['dow'].isin([5,6]).astype(int)

    for lag in [1,7,14,28]:
        ts_df[f'lag_{lag}'] = ts_df['y'].shift(lag)

    for win in [7, 14, 28]:
        ts_df[f'ma_{win}'] = ts_df['y'].shift(1).rolling(win).mean()

    ts_df = ts_df.dropna().reset_index(drop=True)

    print(ts_df.shape)
    display(ts_df.head())

    ###
    try:
        from sklearn.metrics import root_mean_squared_error as rmse_fn
    except ImportError:
        from sklearn.metrics import mean_squared_error
        rmse_fn = lambda y_true, y_pred: mean_squared_error(y_true, y_pred, squared=False)

    cut = int(len(ts_df)*0.8)
    tr, te = ts_df.iloc[:cut], ts_df.iloc[cut:]

    x_tr = tr.drop(columns=['date', 'y']); y_tr = tr['y']
    x_te = te.drop(columns=['date', 'y']); y_te = te['y']

```

```

m = LinearRegression().fit(x_tr, y_tr)
y_hat = m.predict(x_te)

mae = mean_absolute_error(y_te, y_hat)
rmse = rmse_fn(y_te, y_hat)

print('mae:', round(mae,2))
print('rmse:', round(rmse,2))

#%%
def mape(y, yhat):
    y, yhat = np.asarray(y), np.asarray(yhat)
    return np.mean(np.abs((y - yhat) / np.clip(np.abs(y), 1e-8, None))) * 100

def smape(y, yhat):
    y, yhat = np.asarray(y), np.asarray(yhat)
    d = np.abs(y - yhat) / np.clip((np.abs(y) + np.abs(yhat)) / 2, 1e-8, None)
    return np.mean(d) * 100

def bt_lr(df, h=14, start_frac=0.6, feats=('lag_7', 'ma_7', 'ma_14', 'ma_28', 'weekend', 'is_holiday', 'tavg', 'prcp')):
    df = df.sort_values('date').reset_index(drop=True)
    n = len(df)
    start = int(n * start_frac)
    out = []
    last_fold = None

    for t in range(start, n - h):
        tr = df.iloc[:t].dropna(subset=list(feats))
        te = df.iloc[t:t+h]

        x_tr = tr[list(feats)].values
        y_tr = tr['y'].values
        x_te = te[list(feats)].values
        y_te = te['y'].values

        m = LinearRegression().fit(x_tr, y_tr)
        y_hat = m.predict(x_te)
        y_naive = te['lag_7'].values

        mae = mean_absolute_error(y_te, y_hat)
        rmse = np.sqrt(mean_squared_error(y_te, y_hat))

        mape_v = mape(y_te, y_hat)
        smape_v = smape(y_te, y_hat)

        mae_n = mean_absolute_error(y_te, y_naive)
        rmse_n = np.sqrt(mean_squared_error(y_te, y_naive))
        mape_n = mape(y_te, y_naive)
        smape_n = smape(y_te, y_naive)

        out.append({'t0': df.loc[t, 'date'], 'mae': mae, 'rmse': rmse,
                    'mape': mape_v, 'smape': smape_v,
                    'mae_naive': mae_n, 'rmse_naive': rmse_n,
                    'mape_naive': mape_n, 'smape_naive': smape_n})

    last_fold = (te['date'].values, y_te, y_hat, y_naive, mae, rmse)

res = (
    np.mean([r['mae'] for r in out]),
    np.mean([r['rmse'] for r in out]),
    np.mean([r['mape'] for r in out]),
    np.mean([r['smape'] for r in out]),
    np.mean([r['mae_naive'] for r in out]),
    np.mean([r['rmse_naive'] for r in out]),
    np.mean([r['mape_naive'] for r in out]),
    np.mean([r['smape_naive'] for r in out]),
    len(out)
)

dts, y_true, y_hat, y_nv, mae_last, rmse_last = last_fold
plt.figure(figsize=(10, 4))
plt.plot(dts, y_true, label='Actual (y)', color='skyblue', linewidth=2)
plt.plot(dts, y_hat, label='Linear Regression', color='orange', linewidth=2)
plt.plot(dts, y_nv, label='Naive-7', color='limegreen', alpha=0.8)

plt.fill_between(dts, y_hat - mae_last, y_hat + mae_last,

```

```

        color='orange', alpha=0.2, label='±MAE')
plt.fill_between(dts, y_hat - rmse_last, y_hat + rmse_last,
                 color='red', alpha=0.1, label='±RMSE')

plt.legend()
plt.title("Last Fold Forecast with ±MAE and ±RMSE Bands")
plt.xlabel("Date"); plt.ylabel("Revenue")
plt.tight_layout()
plt.show()

print('folds:', res[-1])
print('lr -> mae:', round(res[0],2), 'rmse:', round(res[1],2),
      'mape:', round(res[2],2), 'smape:', round(res[3],2))
print('nv7 -> mae:', round(res[4],2), 'rmse:', round(res[5],2),
      'mape:', round(res[6],2), 'smape:', round(res[7],2))

bt_lr(ts_df, h=14, start_frac=0.6, feats=('lag_7', 'ma_7', 'ma_14', 'ma_28'))

###
def mape(y, yhat):
    y, yhat = np.asarray(y), np.asarray(yhat)
    return np.mean(np.abs((y - yhat) / np.clip(np.abs(y), 1e-8, None))) * 100

def smape(y, yhat):
    y, yhat = np.asarray(y), np.asarray(yhat)
    return np.mean(np.abs(y - yhat) / np.clip((np.abs(y) + np.abs(yhat)) / 2, 1e-8, None)) * 100

def walk_cv(df, h=14, start_frac=0.6, feats=('lag_7', 'ma_7', 'ma_14', 'ma_28')):
    df = df.sort_values('date').reset_index(drop=True)
    n = len(df)
    start = int(n * start_frac)
    folds = []
    for t in range(start, n - h):
        tr = df.iloc[:t].dropna(subset=list(feats))
        te = df.iloc[t:t+h]
        x_tr, y_tr = tr[list(feats)].values, tr['y'].values
        x_te, y_te = te[list(feats)].values, te['y'].values
        folds.append((x_tr, y_tr, x_te, y_te, te['lag_7'].values))
    return folds

def eval_model(model, folds):
    mets = []
    for x_tr, y_tr, x_te, y_te, naive7 in folds:
        m = model.fit(x_tr, y_tr)
        y_hat = m.predict(x_te)
        mets.append({
            'mae': mean_absolute_error(y_te, y_hat),
            'rmse': np.sqrt(mean_squared_error(y_te, y_hat)),
            'mape': mape(y_te, y_hat),
            'smape': smape(y_te, y_hat),
            'mae_nv7': mean_absolute_error(y_te, naive7),
            'rmse_nv7': np.sqrt(mean_squared_error(y_te, naive7)),
            'mape_nv7': mape(y_te, naive7),
            'smape_nv7': smape(y_te, naive7),
        })
    out = pd.DataFrame(mets).mean().to_dict()
    return {k: round(v, 2) for k, v in out.items()}

feats = ('lag_7', 'ma_7', 'ma_14', 'ma_28', 'weekend', 'is_holiday', 'tavg', 'prcp')
folds = walk_cv(ts_df, h=14, start_frac=0.6, feats=feats)

res = []
res.append({'model': 'linreg', **eval_model(LinearRegression(), folds)})
res.append({'model': 'ridge', **eval_model(Ridge(alpha=1.0, random_state=42), folds)})
res.append({'model': 'rf', **eval_model(RandomForestRegressor(
    n_estimators=300, max_depth=8, random_state=42, n_jobs=-1), folds)})

res_df = pd.DataFrame(res).sort_values('rmse').reset_index(drop=True)
print(res_df)

###
feats = ['lag_7', 'ma_7', 'ma_14', 'ma_28', 'weekend', 'is_holiday', 'tavg', 'prcp']

hist_df = (
    ts_df.sort_values('date')
        .reset_index(drop=True)
        .dropna(subset=feats)

```

```

        .copy()
    )
    X = hist_df[feats].values
    y = hist_df['y'].values
    model = Ridge(alpha=1.0, random_state=42).fit(X, y)

    h = 28
    last_date = hist_df['date'].iloc[-1]
    future_dates = pd.date_range(last_date + pd.Timedelta(days=1), periods=h, freq='D')

    fut_df = pd.DataFrame({'date': future_dates})
    fut_df['dow'] = fut_df['date'].dt.dayofweek
    fut_df['woy'] = fut_df['date'].dt.isocalendar().week.astype(int)
    fut_df['month'] = fut_df['date'].dt.month
    fut_df['quarter'] = fut_df['date'].dt.quarter
    fut_df['year'] = fut_df['date'].dt.year
    fut_df['weekend'] = fut_df['dow'].isin([5, 6]).astype(int)

    import holidays
    years = range(int(hist_df['date'].dt.year.min()), int(future_dates[-1].year) + 1)
    us_holidays = holidays.US(years=years)
    holiday_dates = pd.to_datetime(list(us_holidays.keys()))
    fut_df['is_holiday'] = fut_df['date'].isin(holiday_dates).astype(int)

    hist_tavg = dates_df.set_index('date')['tavg'].asfreq('D').ffill() if 'tavg' in dates_df else pd.Series(dtype=float)
    hist_prpcp = dates_df.set_index('date')['prcp'].asfreq('D').fillna(0.0) if 'prcp' in dates_df else pd.Series(dtype=float)
    tavg_seed = float(hist_tavg.tail(28).mean()) if not hist_tavg.empty else 55.0
    prcp_seed = float(hist_prpcp.tail(28).mean()) if not hist_prpcp.empty else 0.0
    fut_df['tavg'] = tavg_seed
    fut_df['prcp'] = prcp_seed

    hist_series = (
        hist_df[['date', 'y']]
        .set_index('date')['y']
        .asfreq('D')
        .ffill()
    )

    def roll_feat(s: pd.Series) -> pd.DataFrame:
        return pd.DataFrame({
            'lag_7': s.shift(7),
            'ma_7': s.rolling(7).mean(),
            'ma_14': s.rolling(14).mean(),
            'ma_28': s.rolling(28).mean()
        })

    work = pd.DataFrame({'y': hist_series.copy()})
    preds = []
    for d in future_dates:
        base = roll_feat(work['y']).iloc[-1][['lag_7', 'ma_7', 'ma_14', 'ma_28']].to_frame().T
        base = base.ffill(axis=1).bfill(axis=1)

        cal = fut_df.loc[fut_df['date'] == d, ['weekend', 'is_holiday', 'tavg', 'prcp']].reset_index(drop=True)

        r = pd.concat([base.reset_index(drop=True), cal], axis=1)[feats]

        yhat = model.predict(r.values).item()
        work.loc[d, 'y'] = yhat
        preds.append(yhat)

    naive7 = hist_series.reindex(future_dates - pd.Timedelta(days=7)).values

    out_df = pd.DataFrame({'date': future_dates, 'yhat_ridge': preds, 'naive7': naive7})

    Path('reports').mkdir(exist_ok=True, parents=True)
    out_df.to_csv('reports/forecast_28d.csv', index=False)
    print('wrote reports/forecast_28d.csv')
    display(out_df.head())

    plt.figure(figsize=(10,4))
    tail_hist = hist_df.tail(90)[['date', 'y']]
    plt.plot(tail_hist['date'], tail_hist['y'], label='y')
    plt.plot(out_df['date'], out_df['yhat_ridge'], label='forecast')
    plt.plot(out_df['date'], out_df['naive7'], label='naive7', alpha=0.6)
    plt.title('28d forecast (ridge)')
    plt.legend(); plt.tight_layout(); plt.show()

```



```

#%%
Path("reports").mkdir(parents=True, exist_ok=True)

last_dt = ts_df['date'].max()
h = len(out_df)

hist_tail = ts_df[ts_df['date'].between(last_dt - pd.Timedelta(days=90), last_dt)][['date', 'y']].copy()

plt.figure(figsize=(10,4))
plt.plot(hist_tail['date'], hist_tail['y'], label='y')
plt.plot(out_df['date'], out_df['yhat_ridge'], label='forecast')
plt.plot(out_df['date'], out_df['naive7'], label='naive7', alpha=0.6)
plt.title(f'{h}d forecast (ridge)')
plt.legend(); plt.tight_layout()
plt.savefig('reports/forecast_plot.png', dpi=150)
plt.close()

fc_sum = out_df['yhat_ridge'].sum()
nv_sum = out_df['naive7'].sum()
hist_28 = ts_df[ts_df['date'].between(last_dt - pd.Timedelta(days=27), last_dt)][['y']].sum()

same_start = (last_dt - pd.Timedelta(days=365 - (h - 1))).normalize()
same_end = same_start + pd.Timedelta(days=h - 1)
yoy_hist = ts_df[ts_df['date'].between(same_start, same_end)][['y']].sum()

summ = pd.DataFrame({
    'metric': ['forecast_sum', 'naive7_sum', 'last_28d_hist', 'same_window_prev_year'],
    'value': [fc_sum, nv_sum, hist_28, yoy_hist]
})
summ.to_csv('reports/forecast_summary.csv', index=False)

with open('reports/forecast_readme.md', 'w') as f:
    f.write(
        f"""# forecast artifacts\n\n
        f"- plot: `reports/forecast_plot.png`\n"
        f"- table: `reports/forecast_28d.csv`\n"
        f"- summary: `reports/forecast_summary.csv`\n\n"
        f"""## quick nums\n
        f"- forecast {h}d sum: {fc_sum:,.0f}\n"
        f"- naive7 {h}d sum: {nv_sum:,.0f}\n"
        f"- last 28d hist sum: {hist_28:,.0f}\n"
        f"- same window prev yr sum: {yoy_hist:,.0f}\n"
    )

print("wrote reports/forecast_plot.png, reports/forecast_summary.csv, reports/forecast_readme.md")

#%%

```